



BIO-4 join

Quail, window GPU 99.9% active



vLLM baseline, window GPU 16.1% active



GPU active Quail executor or vLLM scheduler PyTorch or CUDA API call No recorded CPU operation