


BIO-4 join

Quail

GPU active



CPU operations



0

1

2



3

4

5

seconds

vLLM baseline



0

1

2

3

4

5

seconds

GPU active

Quail executor or vLLM scheduler

PyTorch or CUDA API call