


# BIO-4 join

Quail

GPU active



CPU operations



0

1

2

3

4

5

seconds

vLLM baseline



0

1

2

3

4

5

seconds

GPU active

vLLM scheduler or Quail executor

PyTorch or CUDA API call