


BIO-4 join, vLLM baseline

vLLM baseline

GPU active

Scheduler / execute

PyTorch / CUDA API



GPU active (kernels)

vLLM execute_context

vLLM scheduler or Quail executor

PyTorch or CUDA API call