


BIO-4 join, vLLM baseline

vLLM baseline

GPU active

CPU operations



GPU active

Quail executor or vLLM scheduler

PyTorch or CUDA API call

Other host work