


BIO-3, pipelined vLLM, 480 to 485 seconds after join start

GPU active 1.89 seconds; GPU idle 3.11 seconds



Recorded CPU operations on worker thread 92

vllm.scheduler.schedule



■ PyTorch / CUDA API calls (CPU) ■ vllm.scheduler.* ■ [no recorded CPU operation]