


BIO-3, pipelined vLLM, 480 to 485 seconds after join start

GPU active 1.89 seconds; GPU idle 3.11 seconds



Both panels use the same time axis. CPU children appear below their parent; repeated calls stay separate.

Orange: `vllm.scheduler.*`. Blue: recorded PyTorch and CUDA API calls. Gray: [no recorded CPU operation].

Gray does not mean CPU idle. Open the HTML and click a rectangle to zoom both panels together.

This window is near the join midpoint. Profiling overhead is included; brief intervals can be narrower than a pixel.