



BIO-5 and FEV-11, Qwen34B FP8, sf 0.5, one H100

Quail before Quail now Stock vLLM (operator-at-a-time) SoL estimate


Latency: query time with planning, without model startup (lower is better)



Throughput: every prompt's input tokens over query time (higher is better)



KV regret: fresh tokens minus the token minimum (lower is better)
SoL is 0; stock vLLM's is approximate



Label agreement with Qwen3 32B reference labels (higher is better)

