

BIO-5 and FEV-11, Qwen3 4B FP8, sf 0.5, one H100

Quail before

Quail now

Stock vLLM (operator-at-a-time)

SoL estimate

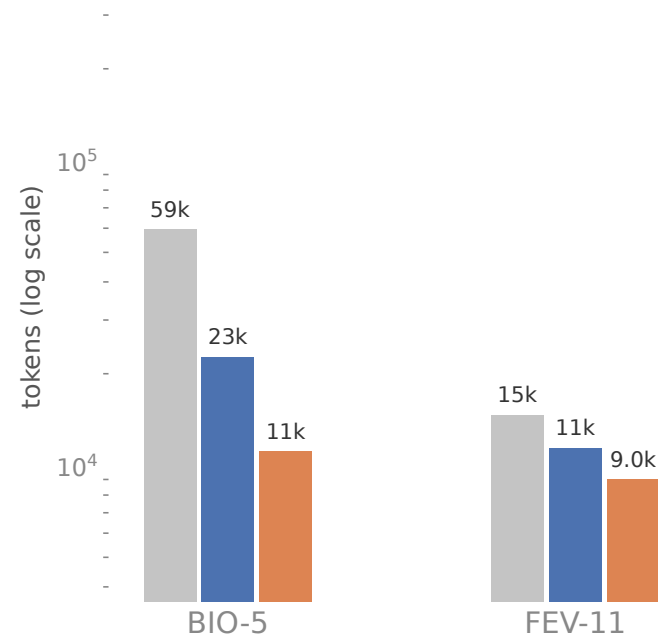
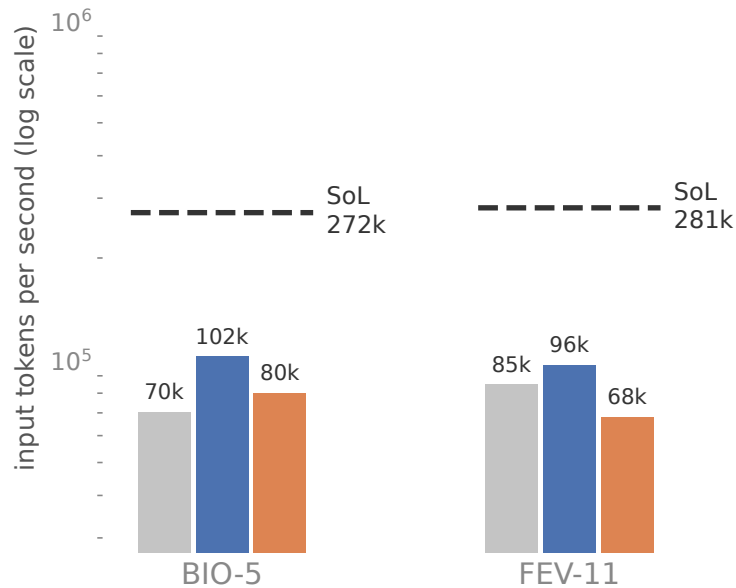
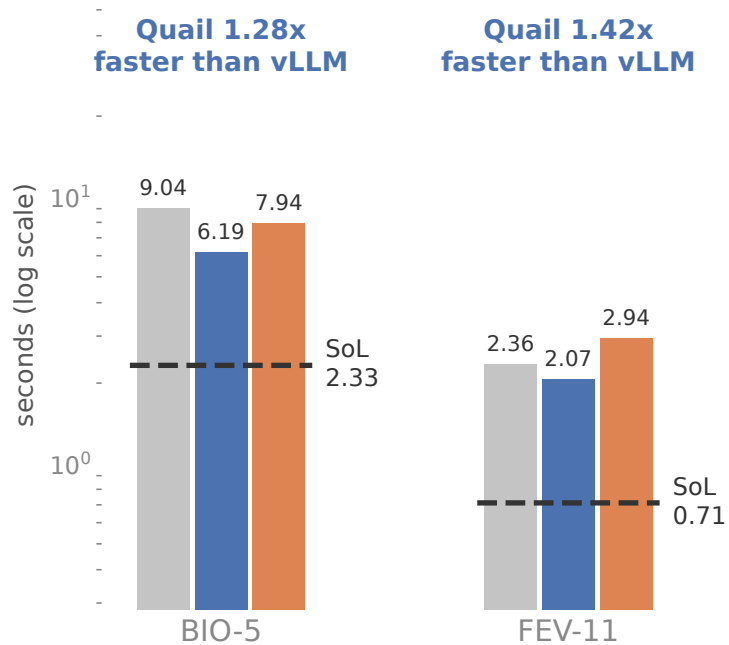
Latency (lower is better)

Throughput (higher is better)

KV regret (lower is better)

Quail 1.28x
faster than vLLM

Quail 1.42x
faster than vLLM



Latency includes planning, not model startup. Throughput is every prompt's input tokens over query time. KV regret is fresh tokens minus the token minimum; SoL's is 0 and stock vLLM's is approximate. SoL: reference-label survivors, unlimited KV.