



DiffusionGemma 26B-A4B FP8, sf 0.5, one H100 per engine

Quail Stock vLLM (operator-at-a-time): not run SoL estimate (distinct prefixes once, unlimited KV, roofline)


Latency: query time with planning, without model startup



Throughput: every prompt's full input length over query time



KV regret: fresh tokens minus the token minimum (SoL is 0 by definition)



Accuracy: label agreement with qwen3-32b-fp8 (no SoL accuracy)

