



QUAIL-B AGENT, Qwen3 4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate


Latency (ratios: vLLM / Quail)




GPU cost per query



GPU cost per million input tokens




Recomputed share of computed tokens



QUAIL-B AGENT, Qwen3 4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Total requested input tokens per second



Raw prompts. References: Qwen3 32B FP8 and dataset annotations. x = not measured.

QUAIL-B AGENT: input documents before filtering

Query	Relation alias (set): input documents
-------	---------------------------------------

AGENT-1	t (agent_traces): 1,772
---------	-------------------------

AGENT-2	t (agent_traces): 1,772
---------	-------------------------

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.
Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.