



# LEP, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate


## Latency (ratios: vLLM / Quail)




## GPU cost per query



## GPU cost per million input tokens




## Recomputed share of computed tokens



# LEP, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail    Pipelined vLLM    SoL estimate

**Total requested input tokens per second**



Reference labels: Qwen3 32B FP8 and dataset annotations.

LEP: input documents before filtering

| Query | Relation alias (set): input documents                  |
|-------|--------------------------------------------------------|
| LEP-1 | d (citation_contexts): 500                             |
| LEP-2 | d (citation_contexts): 500; s (citation_passages): 433 |
| LEP-3 | d (citation_contexts): 500; s (citation_passages): 433 |
| LEP-4 | d (citation_contexts): 500; s (citation_passages): 433 |
| LEP-5 | d (citation_contexts): 500; s (citation_passages): 433 |

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.  
Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.