



LEP, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate


Latency (ratios: vLLM / Quail)




GPU cost per query



GPU cost per million input tokens




Recomputed share of computed tokens



LEP, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Total requested input tokens per second



Reference labels: Qwen3 32B FP8 and dataset annotations.

LEP: input documents before filtering

Query	Relation alias (set): input documents
LEP-1	d (citation_contexts): 500
LEP-2	d (citation_contexts): 500; s (citation_passages): 433
LEP-3	d (citation_contexts): 500; s (citation_passages): 433
LEP-4	d (citation_contexts): 500; s (citation_passages): 433
LEP-5	d (citation_contexts): 500; s (citation_passages): 433

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.
Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.