


QUAIL-B, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Latency




Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Total requested input tokens per second




Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

GPU cost per query




Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

GPU cost per million input tokens




Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B, DiffusionGemma 26B-A4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Recomputed share of computed tokens



Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B: input documents before filtering

Query	Relation alias (set): input documents
IMDB-1	r (reviews): 5,000
IMDB-2	r (reviews): 5,000; a (aspects): 12
IMDB-3	r (reviews): 5,000; a (aspects): 12
IMDB-4	r (reviews): 5,000; a (aspects): 12
IMDB-5	r (reviews): 5,000; a (aspects): 12
IMDB-6	r (reviews): 5,000
IMDB-7	r (reviews): 5,000
IMDB-8	r (reviews): 5,000; a (aspects): 12; a2 (aspects): 12
IMDB-9	r1 (reviews): 5,000; a1 (aspects): 12; r2 (reviews): 5,000; a2 (aspects): 12
IMDB-10	r1 (reviews): 5,000; a1 (aspects): 12; r2 (reviews): 5,000; a2 (aspects): 12
BIO-1	r (reports): 500
BIO-2	r (reports): 500; m (terms): 1,127
BIO-3	r (reports): 500; m (terms): 1,127
FEV-1	c (claims): 500
FEV-2	c (claims): 500; e (evidence): 287
FEV-3	c (claims): 500; e (evidence): 287
FEV-4	c (claims): 500; e (evidence): 287
FEV-5	c (claims): 500; e (evidence): 287
FEV-6	c (claims): 500; e (evidence): 287
FEV-7	c (claims): 500; e (evidence): 287; e2 (evidence): 287
FEV-8	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-9	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-10	c (claims): 500; e (evidence): 287
LEP-1	d (citation_contexts): 500
LEP-2	d (citation_contexts): 500; s (citation_passages): 433
LEP-3	d (citation_contexts): 500; s (citation_passages): 433
LEP-4	d (citation_contexts): 500; s (citation_passages): 433
LEP-5	d (citation_contexts): 500; s (citation_passages): 433
AGENT-1	t (agent_traces): 1,772
AGENT-2	t (agent_traces): 1,772

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.
Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.