



QUAIL-B FEV, Qwen3 4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate


Latency (ratios: vLLM / Quail)




Fresh input tokens



Recomputed KV tokens



Answer agreement with reference labels




Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B FEV, Qwen3 4B FP8, sf=0.1, one H100

Quail Pipelined vLLM SoL estimate

Total requested input tokens per second



Reference labels: Qwen3 32B FP8 and dataset annotations.

QUAIL-B FEV: input documents before filtering

Query	Relation alias (set): input documents
FEV-1	c (claims): 500
FEV-2	c (claims): 500; e (evidence): 287
FEV-3	c (claims): 500; e (evidence): 287
FEV-4	c (claims): 500; e (evidence): 287
FEV-5	c (claims): 500; e (evidence): 287
FEV-6	c (claims): 500; e (evidence): 287
FEV-7	c (claims): 500; e (evidence): 287; e2 (evidence): 287
FEV-8	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-9	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-10	c (claims): 500; e (evidence): 287

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.

Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.