



QUAIL-B FEV, Qwen3 4B FP8, sf=0.1, one H100

Quail Pipelined stock vLLM SoL estimate


Latency (ratios: vLLM / Quail)




Fresh input tokens computed



Recomputed KV tokens



Answer agreement



Raw prompts. References: Qwen3 32B FP8 and dataset annotations. x = not measured.

QUAIL-B FEV: input documents before filtering

Query	Relation alias (set): input documents
FEV-1	c (claims): 500
FEV-2	c (claims): 500; e (evidence): 287
FEV-3	c (claims): 500; e (evidence): 287
FEV-4	c (claims): 500; e (evidence): 287
FEV-5	c (claims): 500; e (evidence): 287
FEV-6	c (claims): 500; e (evidence): 287
FEV-7	c (claims): 500; e (evidence): 287; e2 (evidence): 287
FEV-8	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-9	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-10	c (claims): 500; e (evidence): 287

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.

Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.