



QUAIL-B LEP, Qwen3 4B FP8, sf=0.1, one H100

Quail Stock vLLM Pipelined vLLM Pipelined SGLang SoL estimate


Latency




Fresh input tokens



Recomputed KV tokens



Answer agreement with Qwen3 32B



SoL estimates ideal work with unlimited prefix KV reuse across requests and reference-label survivors. It has no measured accuracy.
Fresh tokens include recomputed KV. A dash marks zero. Stock vLLM uses operator-at-a-time submission.
SGLang measurements use the earlier adapter.

QUAIL-B LEP: input documents before filtering

Query	Relation alias (set): input documents
LEP-1	d (citation_contexts): 500
LEP-2	d (citation_contexts): 500; s (citation_passages): 433
LEP-3	d (citation_contexts): 500; s (citation_passages): 433
LEP-4	d (citation_contexts): 500; s (citation_passages): 433
LEP-5	d (citation_contexts): 500; s (citation_passages): 433
LEP-6	d (citation_contexts): 500; s (citation_passages): 433
LEP-7	d (citation_contexts): 500; s (citation_passages): 433
LEP-8	d (citation_contexts): 500

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.
Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.