


QUAIL-B, Qwen3 4B FP8, sf=0.1, one H100




SoL estimates ideal work with unlimited prefix KV reuse across requests and reference-label survivors. It has no measured accuracy.

Fresh tokens include recomputed KV. A dash marks zero. Stock vLLM uses operator-at-a-time submission.

FEV-9 uses the revised SGLang adapter; other SGLang results use the earlier adapter.

QUAIL-B, Qwen3 4B FP8, sf=0.1, one H100




SoL estimates ideal work with unlimited prefix KV reuse across requests and reference-label survivors. It has no measured accuracy.

Fresh tokens include recomputed KV. A dash marks zero. Stock vLLM uses operator-at-a-time submission.

FEV-9 uses the revised SGLang adapter; other SGLang results use the earlier adapter.

QUAIL-B, Qwen3 4B FP8, sf=0.1, one H100



SoL estimates ideal work with unlimited prefix KV reuse across requests and reference-label survivors. It has no measured accuracy.


Fresh tokens include recomputed KV. A dash marks zero. Stock vLLM uses operator-at-a-time submission.

FEV-9 uses the revised SGLang adapter; other SGLang results use the earlier adapter.

QUAIL-B, Qwen3 4B FP8, sf=0.1, one H100

Quail Stock vLLM Pipelined vLLM Pipelined SGLang

Answer agreement with Qwen3 32B



SoL estimates ideal work with unlimited prefix KV reuse across requests and reference-label survivors. It has no measured accuracy.

Fresh tokens include recomputed KV. A dash marks zero. Stock vLLM uses operator-at-a-time submission.

FEV-9 uses the revised SGLang adapter; other SGLang results use the earlier adapter.

QUAIL-B: input documents before filtering

Query	Relation alias (set): input documents
IMDB-1	r (reviews): 5,000
IMDB-2	r (reviews): 5,000; a (aspects): 12
IMDB-3	r (reviews): 5,000; a (aspects): 12
IMDB-4	r (reviews): 5,000; a (aspects): 12
IMDB-5	r (reviews): 5,000; a (aspects): 12
IMDB-6	r (reviews): 5,000
IMDB-7	r (reviews): 5,000
IMDB-8	r (reviews): 5,000; a (aspects): 12; a2 (aspects): 12
IMDB-9	r1 (reviews): 5,000; a1 (aspects): 12; r2 (reviews): 5,000; a2 (aspects): 12
IMDB-10	r1 (reviews): 5,000; a1 (aspects): 12; r2 (reviews): 5,000; a2 (aspects): 12
BIO-1	r (reports): 500
BIO-2	r (reports): 500; m (terms): 1,127
BIO-3	r (reports): 500; m (terms): 1,127
FEV-1	c (claims): 500
FEV-2	c (claims): 500; e (evidence): 287
FEV-3	c (claims): 500; e (evidence): 287
FEV-4	c (claims): 500; e (evidence): 287
FEV-5	c (claims): 500; e (evidence): 287
FEV-6	c (claims): 500; e (evidence): 287
FEV-7	c (claims): 500; e (evidence): 287; e2 (evidence): 287
FEV-8	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
FEV-9	c1 (claims): 500; e1 (evidence): 287; c2 (claims): 500; e2 (evidence): 287
LEP-1	d (citation_contexts): 500
LEP-2	d (citation_contexts): 500; s (citation_passages): 433
LEP-3	d (citation_contexts): 500; s (citation_passages): 433
LEP-4	d (citation_contexts): 500; s (citation_passages): 433
LEP-5	d (citation_contexts): 500; s (citation_passages): 433
LEP-6	d (citation_contexts): 500; s (citation_passages): 433
LEP-7	d (citation_contexts): 500; s (citation_passages): 433
LEP-8	d (citation_contexts): 500
AGENT-1	t (agent_traces): 1,772
AGENT-2	t (agent_traces): 1,772

Each alias lists its full input before filters. Repeated aliases can refer to the same underlying set.

Counts come from the saved corpus manifest. SoL uses the same inputs and the saved reference labels for survivors.